

1. The reason for selecting Continual Learning

Deep Learning, a prevailing class of machine learning algorithms, requires vast labelled samples to train models so that a specific isolated task can be tackled.

This kind of training is classified as offline training and supervised learning, and machine intelligence endowed with general knowledge and exceptional perception scarcely develops through this sort of learning.

Also, it is arduous to acquire enormous labelled training samples from intricate problems in the real world, which makes scholars shift their fields of research to unsupervised learning gradually. Convoluted problems are addressed by applying the features of unsupervised learning.

Further, future artificial intelligence systems usually demand outstanding potential for perception and require adaptation to confront external problems through learning from its surroundings incessantly. Hence, some scholars think that future artificial intelligence systems should behave like humans not only able to constantly learn from and adjust to the exterior world but also autonomously develop more complex skills and understanding. Thus, Continual Learning is an ideal preference for unsupervised learning.

Continual Learning is not only able to deal with high-dimensional and real-time data but also continually update models which adapt to and resolve problems of new situations when surroundings change over time, aka new usable data inputted.

Currently, obtainable frameworks and algorithms for Continual Learning exist. I select Caffe as my Deep Learning framework and opt for ResNet featuring continual learning as my algorithm.

ResNet, categorized into an unsupervised algorithm, can regularly learn from and adapt to an environment. Also, ResNet autonomously evolves more complicated skills and knowledge. A thesis, "Deep Residual Learning for Image Recognition", states that ResNet can gain fine experiences from low levels and transfer these experiences to high levels. Learning is endowed with memorization to avoid catastrophic forgetting, vanishing gradients, exploding gradients, and degradation problem.

ResNet exploits an extra shortcut path to interface with the next layer or next layers. Also, the sizes of the layers have to be the same size so that the outcomes trained by low levels can be transferred into high levels. Thus, ResNet can combine artlessly low/mid/high-level attributes and classifiers, which described as an end-to-end multi-layer fashion in the thesis, “Deep Residual Learning for Image Recognition”.


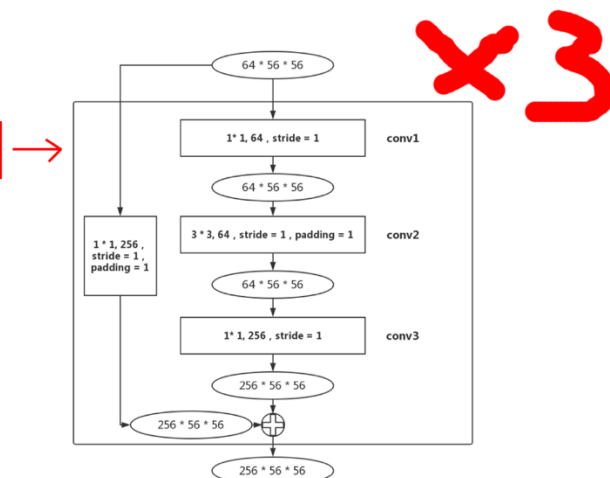


Figure 2. Residual learning: a building block.

Take ResNet-50 as an example to elucidate the mechanism of a ResNet-50 architecture.

- (1) The layers before the main blocks of ResNet-50 (the lowest layers) include conv1 (7×7 convolution layer) and conv2_x (3×3 max pooling layer).
- (2) Layer 1 is Conv2_x (3×3×64 as a Basicblock). One convolution layer is attached to an input channel and an output channel. The kernel size of an output channel is enlarged 4 times that of an input channel. A bottleneck consists of an input channel, an output channel and a Basicblock.

layer name	output size	50-layer
conv1	112×112	7×7, 64, stride 2 3×3 max pool, stride 2
conv2_x	56×56	<div style="border: 1px solid red; padding: 2px; display: inline-block;"> 1×1, 64 3×3, 64 1×1, 256 </div> ×3
conv3_x	28×28	1×1, 128 3×3, 128 1×1, 512
conv4_x	14×14	1×1, 256 3×3, 256 1×1, 1024
conv5_x	7×7	1×1, 512 3×3, 512 1×1, 2048
	1×1	average pool, 1000-d fc, softmax



- (3) Excepting the number of channels expanded and the kernel size of output channels is unlike, the constitution of Layer2, Layer3 and Layer4 are the same as that of Layer1.

2. The procedures for executing my program with BVLC/ResNet architectures

Step 1: Set up a programming environment:

I followed the instructions from the file "README.md" on the website (vlomonaco.github.io/core50/index.html) designated by the professor to establish the programming environment. Then, I downloaded code from the URL (github.com/vlomonaco/core50) and downloaded data by executing the bash file "fetch_data_and_setup.sh".

Step 2: Configure and execute the program:

I comprehended the entire architecture of the program from the folders "core" and "confs" in the project. Next, I modified the files in the folder "core" to eliminate unnecessary code such as code for different classification strategies and Lightning Memory-Mapped Database. After that, I executed the BVLC model with the NC configurations to grasp the mechanism of the program. To overcome the abysmal load speed of images, I opted for a RAM disk to hasten to load images.

Step 3: Choose ResNet-10/ResNet-50 to reproduce the experiments:

I downloaded ResNet-10/ResNet-50 models and respective configuration files from @cvjena's Github (github.com/cvjena/cnn-models/tree/master/ResNet_preact), altered the configurations and reproduced the experiments.

Step 4: Benchmarks:

Benchmarks (9-batches/1-run) for classification at object level (50 classes) applying BVLC, ResNet-10 and ResNet-50 are as follows:

BVLC:

```
Batch: 8, name: train_batch_08 (size 119894), Accuracy: 0.6568444444444445
[Train-net] avg. new weights: 2.0324189e-05
[Train-net] avg. new biases: -2.9343647e-07
[Train-net] avg. other weights: nan
[Train-net] avg. other biases: nan
[Train-net] tot weights avg. : 2.0324189e-05
[Test-net] avg. new weights: 2.0324189e-05
[Test-net] avg. new biases: -2.9343647e-07
[Test-net] avg. other weights: nan
[Test-net] avg. other biases: nan
Acc. per class:
0: 0.747      1: 0.251      2: 0.64       3: 0.284      4: 0.492
5: 0.697      6: 0.625      7: 0.43       8: 0.263      9: 0.548
10: 0.589     11: 0.579     12: 0.387     13: 0.687     14: 0.853
15: 0.641     16: 0.201     17: 0.594     18: 0.805     19: 0.329
20: 0.743     21: 0.796     22: 0.835     23: 0.932     24: 0.586
25: 0.576     26: 0.385     27: 0.732     28: 0.586     29: 0.31
30: 0.982     31: 0.909     32: 0.914     33: 0.98       34: 0.933
35: 0.835     36: 0.925     37: 0.793     38: 0.571     39: 0.751
40: 0.868     41: 0.468     42: 0.72       43: 0.726     44: 0.863
45: 0.729     46: 0.731     47: 0.697     48: 0.602     49: 0.72
```

ResNet-10:

```
Batch: 8, name: train_batch_08 (size 119894), Accuracy: 0.6843555555555556
[Train-net] avg. new weights: 0.009121395
[Train-net] avg. new biases: 0.021162495
[Train-net] avg. other weights: nan
[Train-net] avg. other biases: nan
[Train-net] tot weights avg. : 0.009121395
[Test-net] avg. new weights: 0.009121395
[Test-net] avg. new biases: 0.021162495
[Test-net] avg. other weights: nan
[Test-net] avg. other biases: nan
Acc. per class:
0: 0.809      1: 0.469      2: 0.641      3: 0.347      4: 0.634
5: 0.648      6: 0.629      7: 0.274      8: 0.33       9: 0.472
10: 0.711     11: 0.51      12: 0.578     13: 0.64      14: 0.952
15: 0.58      16: 0.404     17: 0.699     18: 0.887     19: 0.489
20: 0.759     21: 0.802     22: 0.76      23: 0.851     24: 0.587
25: 0.879     26: 0.667     27: 0.77      28: 0.623     29: 0.383
30: 0.964     31: 0.97      32: 0.928     33: 0.964     34: 0.944
35: 0.857     36: 0.832     37: 0.71      38: 0.628     39: 0.795
40: 0.89      41: 0.508     42: 0.77      43: 0.651     44: 0.734
45: 0.635     46: 0.774     47: 0.753     48: 0.497     49: 0.627
```

ResNet-50:

```
Batch: 8, name: train_batch_08 (size 119894), Accuracy: 0.8215333333333333
[Train-net] avg. new weights: 0.0030594722
[Train-net] avg. new biases: 0.011786158
[Train-net] avg. other weights: nan
[Train-net] avg. other biases: nan
[Train-net] tot weights avg. : 0.0030594722
[Test-net] avg. new weights: 0.0030594722
[Test-net] avg. new biases: 0.011786158
[Test-net] avg. other weights: nan
[Test-net] avg. other biases: nan
Acc. per class:
0: 0.927      1: 0.687      2: 0.723      3: 0.489      4: 0.827
5: 0.886      6: 0.572      7: 0.688      8: 0.621      9: 0.713
10: 0.942     11: 0.733     12: 0.459     13: 0.849     14: 0.938
15: 0.782     16: 0.669     17: 0.855     18: 0.956     19: 0.893
20: 0.861     21: 0.879     22: 0.897     23: 0.963     24: 0.829
25: 0.782     26: 0.813     27: 0.887     28: 0.717     29: 0.504
30: 0.984     31: 0.993     32: 0.961     33: 0.997     34: 0.978
35: 0.969     36: 0.909     37: 0.859     38: 0.818     39: 0.844
40: 0.906     41: 0.812     42: 0.809     43: 0.796     44: 0.92
45: 0.845     46: 0.879     47: 0.892     48: 0.77      49: 0.798
```

From the benchmarks (9-batches/1-run) for classification at object level (50 classes), the performance of ResNet-50 prevails over ResNet-10 and BVLC.